

Posterior Sampling with Proximal MCMC

Louis Christie
lgc26@cam.ac.uk

March 5, 2021

1 Introduction

In many image settings we are interested in the inverse problem of finding $x \in \mathbb{X}$ when we see $y \sim N(Ax, V)$ for example. One way to do this is to find a posterior distribution

$$\pi_{X|Y}(x | y) \propto \pi_{Y|X}(y | x) \times \pi_X(x) = N(y | Ax, V) \times \pi_X(x) \quad (1.1)$$

This requires us to specify a prior $\pi_X(x)$, which is often done by training a neural network or other model on a set of accurate images.

Once we've specified the prior and likelihood, we can estimate x from the posterior. One way is to find the Maximum A Posteriori (MAP) estimate $\hat{x}^{MAP} = \arg \max_{u \in \mathbb{X}} \pi_{X|Y}(u | y)$. Another is to find the posterior expectation $\hat{x} = \mathbb{E}(X | y)$. Since it may not be possible to analytically find this expectation, we turn to Monte-Carlo integration:

$$\mathbb{E}(X | Y) \approx \frac{1}{m} \sum_{i=1}^m X_i \quad \text{where } X_i \stackrel{iid}{\sim} \pi_{X|Y} \quad (1.2)$$

To do this we need to sample from the posterior, which we do by designing a Markov Chain with the target posterior as its stationary distribution. In our case we can use the MYULA chain

Proposition 1.1. *If $-\log \pi_{X|Y} = f + g$ where f is convex, proper, lower semicontinuous and has L_f -Lipschitz gradient and g is convex, proper and lower semicontinuous, then the chain with transitions:*

$$X_{n+1} = X_n - \delta \nabla f(X_n) - \frac{\delta}{\lambda} (X_n - \text{Prox}_g^\lambda(X_n)) + \sqrt{2\delta} Z_{n+1} \quad (\text{MYULA})$$

(where $Z_i \stackrel{iid}{\sim} N(0, 1)$ and $\delta < 2\lambda(L_f\lambda + 1)^{-1}$) converges exponentially fast to $\pi_{X|Y}$ as $n \rightarrow \infty$.

From equation 1.1 we see that for some normalisation constant C :

$$-\log \pi_{X|Y} = -\log C - \log \pi_{Y|X} - \log \pi_X \quad (1.3)$$

So we can set $f = -\log C - \log \pi_{Y|X}$ and $g = -\log \pi_X$ to use MYULA. This requires computing $\text{Prox}_g^\lambda(X_n) = \arg \min_{u \in \mathbb{X}} (g(u) + (2\lambda)^{-1} \|x - u\|_{\mathbb{X}}^2)$ at each step. This is a strongly convex function so does have a unique minimiser, but we would like to compute it without a descent style algorithm. The question becomes can we learn a prior π_X such that this proximal operator can be computed analytically?

Since g is sub-differentiable, we can define see that the proximal operator can be rewritten as an inverse image:

$$\text{Prox}_g^\lambda(x) = (I_{\mathbb{X}} + \lambda \partial g)^{-1}(\{x\}) \quad (1.4)$$

This can be seen as if $u \in (I_{\mathbb{X}} + \lambda \partial g)(x)$ then $0 \in \lambda \partial g(u) + u - x$ so u is a minimiser of $g(u) + \|x - u\|_{\mathbb{X}}^2 / (2\lambda)$. Since this is strongly convex, the inverse image is a singleton.

[cite Subhadip](#) have shown that they can learn g as an **input convex neural network** (ICNN) by using nodes of the form:

$$z_{i+1} = \phi_i(B_i(z_i) + W_i(x_i) + b_i)$$

If the weights B_i are all non-negative and the ϕ_i are each convex and monotone then the learned network is also convex.

If we use a smooth activation function ϕ_i , this is also guaranteed to be smooth. This means that we can ignore the proximal step in MYULA entirely - we can just use the smooth posterior to sample easily.

2 Gaussian Example

Suppose that A is the identity operator, so $y \mid x \sim N(x, \Sigma_\epsilon)$. If we take a Gaussian prior on X , so $\pi_X(x) = N(x \mid 0, \Sigma_X)$, then we have a known posterior:

$$\pi_{X|Y}(x \mid y) \propto N(x \mid y, \Sigma_\epsilon) \times N(x \mid 0, \Sigma_X) \quad (2.1)$$

$$\propto N(x \mid S(\Sigma_\epsilon^{-1}y), S) \quad (2.2)$$

where $S^{-1} = \Sigma_\epsilon^{-1} + \Sigma_X^{-1}$.

Example 2.1. In Willem's example, we have $y = (1, 1)$, $\Sigma_\epsilon = \frac{1}{2}I_2$, and $\Sigma_X = I_2$. Thus we have the posterior

$$\pi_{X|Y}(x \mid (1, 1)^T) = N\left(x \mid \begin{pmatrix} \frac{3}{2} \\ \frac{3}{2} \end{pmatrix}, \begin{pmatrix} \frac{1}{3} & 0 \\ 0 & \frac{1}{3} \end{pmatrix}\right) \quad (2.3)$$

Sampling 4000 steps of MYULA with 1000 to burn in gives a sample with a KDE given in figure 2.1.

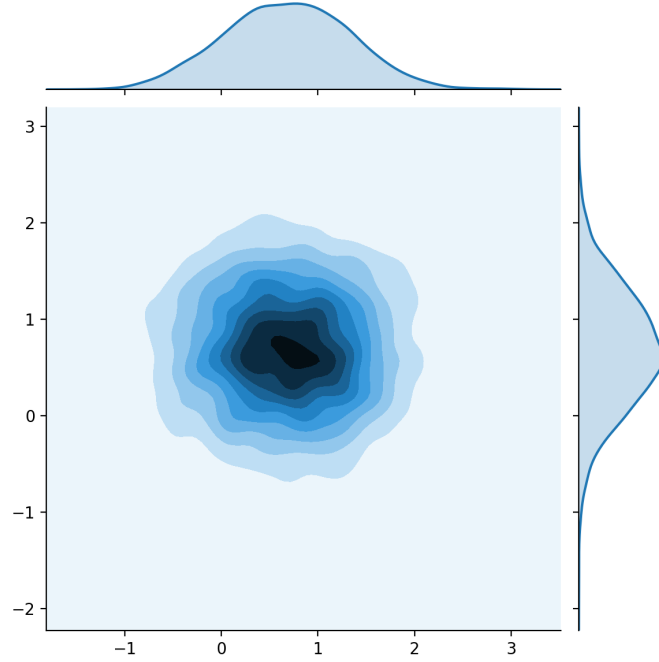


Figure 1: Posterior Sample KDE from MYULA. Sample mean at $(0.682, 0.681)^T$ with marginal variances of 0.42 and 0.38

3 Lasso Prior

In the same context as above we have the likelihood $N(y \mid Ax, \Sigma_\epsilon)$, but now we have $-\log \pi_X(x) = g(x) = \|x\|_1$. Thus the prox of g is given by:

$$\text{Prox}_g^\lambda(x) = \arg \min_{u \in \mathbb{R}^d} \frac{1}{2} \|u - x\|_2^2 + \lambda \|u\|_1 \quad (3.1)$$

As the objective is separable componentwise; we get:

$$\text{Prox}_g^\lambda(x)_i = \arg \min_{u_i \in \mathbb{R}} \frac{1}{2} (u_i - x_i)^2 + \lambda |u_i| \quad (3.2)$$

$$= \{u_i \in \mathbb{R} : 0 \in \partial \frac{1}{2} (u_i - x_i)^2 + \lambda |u_i|\} \quad (3.3)$$

$$= \{u_i \in \mathbb{R} : 0 \in (u_i - x_i) + \lambda \text{sign}(u_i) \mathbf{1}_{u_i \neq 0} + \lambda \mathbf{1}_{u_i = 0} [-1, 1]\} \quad (3.4)$$

$$= \text{sign}(x_i) \max(0, |x_i| - \lambda) \quad (3.5)$$

This is called the **soft-thresholding operator**.




Figure 2: Posterior Sample KDE from MYULA. Sample mean at $(0.639, -0.019)^T$ with marginal variances of 0.44 and 0.34

4 Total Variation Prior

5 Smooth Input Convex Neural Networks

The nodes of ICNNs are of the form:

$$z_{i+1} = \phi_i(B_i(z_i) + W_i(x_i) + b_i)$$

so if we use a smooth activation function $\phi_i = \text{softmax}$ we will have a smooth g , allowing us to ignore the prox step entirely.