

# Contents

<b>1</b>	<b>Learned Convex Prior Sampling</b>	<b>2</b>
1.1	Core papers and ideas . . . . .	2
1.2	Ideas for the project . . . . .	3
<b>2</b>	<b>Proximal Map Learning</b>	<b>4</b>
2.1	Main Idea . . . . .	4
2.2	First Steps: a Toy Example . . . . .	4

# Todo list

---

# Section 1: Learned Convex Prior Sampling

## 1.1 Core papers and ideas

- Exponential convergence of Langevin distributions and their discrete approximations [4]
    - Propose ULA
      - \* Thm 2.1: if our convex regularizer is smaller than  $a\|x\|^2 + b$  for large  $\|x\|$  we can expect convergence. We don't know the rate at this point
      - \* Thm 2.3 does give a criterion for geometric convergence
      - \* Thm 2.4 gives criterion for when we don't get the rate (or convergence whatsoever?)
    - And propose MALA and give similar convergence results with rates
    - At the end of this paper it is not very clear whether MALA is much better than ULA: both seem to converge geometrically under the same conditions
      - \* According to [5]: the difference between ULA and MALA seems to be a bias for the former
        - so we don't converge to the correct distribution
        - only for very small step sizes, but then the method becomes infeasible
  - Proximal markov chain monte carlo algorithms [3]
    - Instead of doing a forward Euler gradient step, we can also do backward Euler, i.e., a proximal step
    - Propose P-ULA and P-MALA
      - \* Thm 3.1: we get similar result as Thm 2.3 in [4] for P-ULA
      - \* Thm 3.4: geometric convergence for P-MALA
    - Use FBS to get a decent expression for the prox of the whole functional
      - \* No proof here how accurate the FBS approximation is, but empirically very accurate (also for high dimensional models)
      - \* Later looked into again [?] and is now known as MYULA (and MY-MALA?)
    - Applying the models still relies on having a good approximation of the proximal map
      - \* Here they use TV and use [1] for the TV algorithm
      - \* Duality based and not very generalizable to other functionals
  - Accelerating proximal Markov chain Monte Carlo by using an explicit stabilised method [5]
    - Speed-up, but no theoretical contributions
-

- Learned convex regularizers for inverse problems [2]
  - learning regularizers

## 1.2 Ideas for the project

Main three tasks

- Set up MYULA
- Set up Regulariser learner
- Set up Prox solver

Further to start with

- Get MYMALA/MYULA to work on something we know the distribution of
  - 2D Gaussian and tikhonov regularization
    - \* also implement ULA/MALA for comparison
  -

---

## Section 2: Proximal Map Learning

### 2.1 Main Idea

We want to be able to do the Markov step

$$X_{n+1} = X_n - \delta \nabla f(X_n) - \frac{\delta}{\lambda} \left( X_n - \text{prox}_g^\lambda(X_n) \right) + \sqrt{2\delta} Z_{n+1} \quad (1)$$

where  $f$  our data fidelity term and  $g = \mathcal{R}_\theta$  our learned regularizer.

Note that in the construction of the Markov chain we used

$$\begin{aligned} \nabla \log \pi^\lambda &= -\nabla f(x) - \nabla g^\lambda(x) \\ &= -\nabla f(x) - \frac{1}{\lambda} \left( x - \text{prox}_g^\lambda(x) \right) \end{aligned} \quad (2)$$

Instead of learning the prox directly, we can also try to learn  $g^\lambda$ . The immediate problem here is:

- the function has scalar (1D) output: dimensionality-wise efficient representation
- the function is differentiable
- the function is convex!

So we can use the input-convex neural networks (ICNNs) framework

Then for the resulting NN  $\mathcal{H}_\phi$  we can either plug in its gradient right away or we can compute the proximal map explicitly through

$$\text{prox}_g^\lambda(x) = x - \lambda \nabla \mathcal{H}_\phi(x) \quad (3)$$

### 2.2 First Steps: a Toy Example

We will first focus on an example of something we already know the answer of:  $g(x) = \alpha \|x\|_1$ . There we have for  $x \in \mathbb{R}^N$

$$g^\lambda(x) = \inf_y \left\{ \frac{1}{2} \|x - y\|_2^2 + \lambda \alpha \|y\|_1 \right\} = \sum_i^N S_{\lambda\alpha}(x_i) \quad (4)$$

where

$$S_{\lambda\alpha}(z) = \begin{cases} \frac{1}{2\lambda\alpha} z^2, & |z| \leq \lambda\alpha \\ |z| - \frac{\lambda\alpha}{2}, & |z| > \lambda\alpha \end{cases} \quad (5)$$

works coordinate wise, i.e.,  $x = (x_i)_i$ .

---

### Data generation

So we want to train a network  $\mathcal{H}_\phi$  that approximates  $g^\lambda$ . We will initially use a supervised approach with data  $\{x_j, y_j\}$  where  $x_j \sim \pi_X$  and  $y_j = g^\lambda(x_j)$ <sup>1</sup>

- Fix  $\lambda$  (we can also do without this assumption, but initially this will be easier for data generation and training)
- At this point it is not entirely clear how to sample from  $\pi_X$  in order to generate our  $x_j$
- I guess normally, we do have some  $x_j$ 's already which we used to train our  $\mathcal{R}_\theta$ . We should reuse them then
- In that sense, it would make sense to sample from the *prior distribution*.
  - Then, for our toy example we will use  $\pi_X = \text{Laplace}(0, \frac{2}{\alpha^2})^N$

$$x_j \sim \frac{1}{2\alpha} e^{-\alpha\|x\|_1} dx \quad (6)$$

corresponding to our prior.

- We should note that we have typically trained a regularizer  $\mathcal{R}_\theta$  and not  $\alpha\mathcal{R}_\theta$ .
  - \* So maybe we shouldn't use  $\alpha$  in the distribution for the toy example either
  - \* Or we should just set  $\alpha = 1$  and assume that our regularizer is just fine as it is.

### The Learning problem

We need to formulate a loss function to define the learning problem. We will initially choose plain  $\ell^2$  loss, i.e.,

$$\phi^* = \arg \min_{\phi} \left( \mathbb{E}_{\pi_X} \left[ \left( \mathcal{H}_\phi(X) - g^\lambda(X) \right)^2 \right] \right) \quad (7)$$

and in order to do this, we will minimize the empirical risk

$$\phi^* = \arg \min_{\phi} \frac{1}{N} \sum_j^N \|\mathcal{H}_\phi(x_j) - y_j\|^2. \quad (8)$$

Additionally we will potentially use  $\ell^2$  regularization on the weights to prevent overfitting.

### Network Architecture

ICNN

---

<sup>1</sup>For a general  $\mathcal{R}_\theta$ , we just need to solve the problem here a couple of times to generate data

## References

- [1] Antonin Chambolle. An algorithm for total variation minimization and applications. *Journal of Mathematical imaging and vision*, 20(1):89–97, 2004.
- [2] Subhadip Mukherjee, Sören Dittmer, Zakhar Shumaylov, Sebastian Lunz, Ozan Öktem, and Carola-Bibiane Schönlieb. Learned convex regularizers for inverse problems. *arXiv preprint arXiv:2008.02839*, 2020.
- [3] Marcelo Pereyra. Proximal markov chain monte carlo algorithms. *Statistics and Computing*, 26(4):745–760, 2016.
- [4] Gareth O Roberts, Richard L Tweedie, et al. Exponential convergence of langevin distributions and their discrete approximations. *Bernoulli*, 2(4):341–363, 1996.
- [5] Luis Vargas, Marcelo Pereyra, and Konstantinos C Zygalakis. Accelerating proximal markov chain monte carlo by using an explicit stabilised method. *arXiv preprint arXiv:1908.08845*, 2019.